

Collaborative Filtering Overview 1

BY GUOKAN SHANG


Research Engineer @ Labs, Linagora &
PhD Candidate @ DaSciM, LIX

Email: `guokan.shang@hotmail.com`

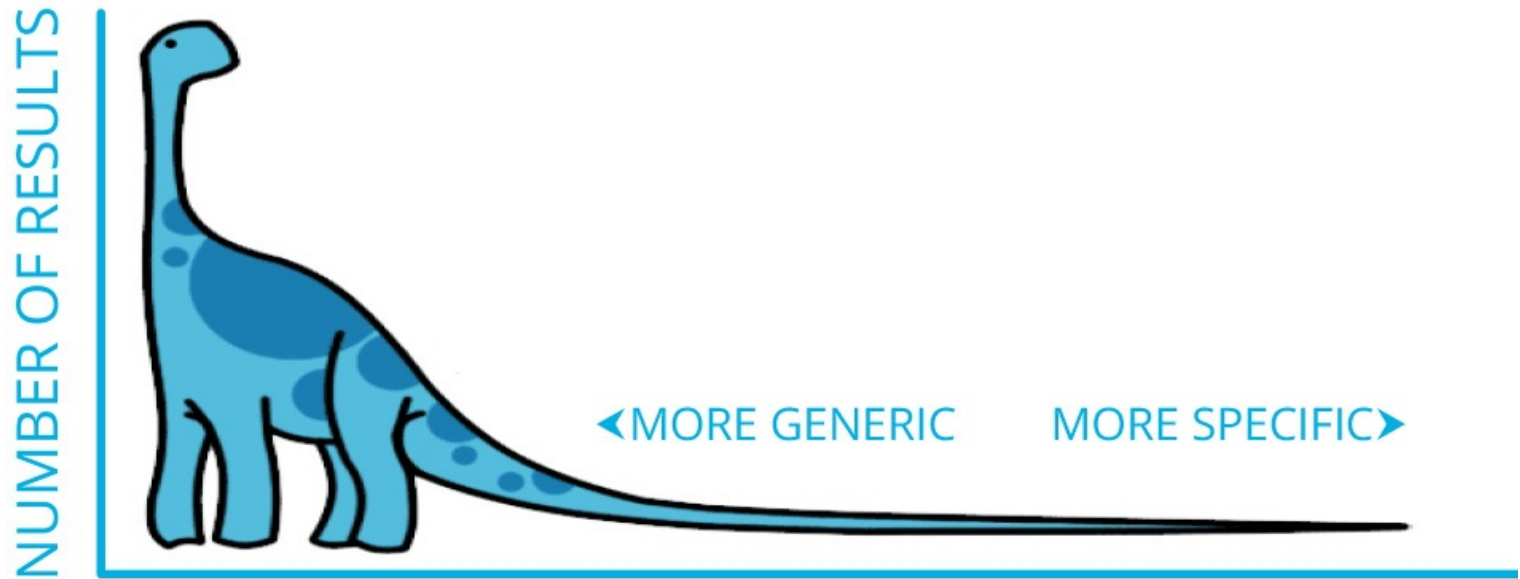
Web: `http://www.lix.polytechnique.fr/Labo/Guokan.Shang/`

July 27, 2017

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

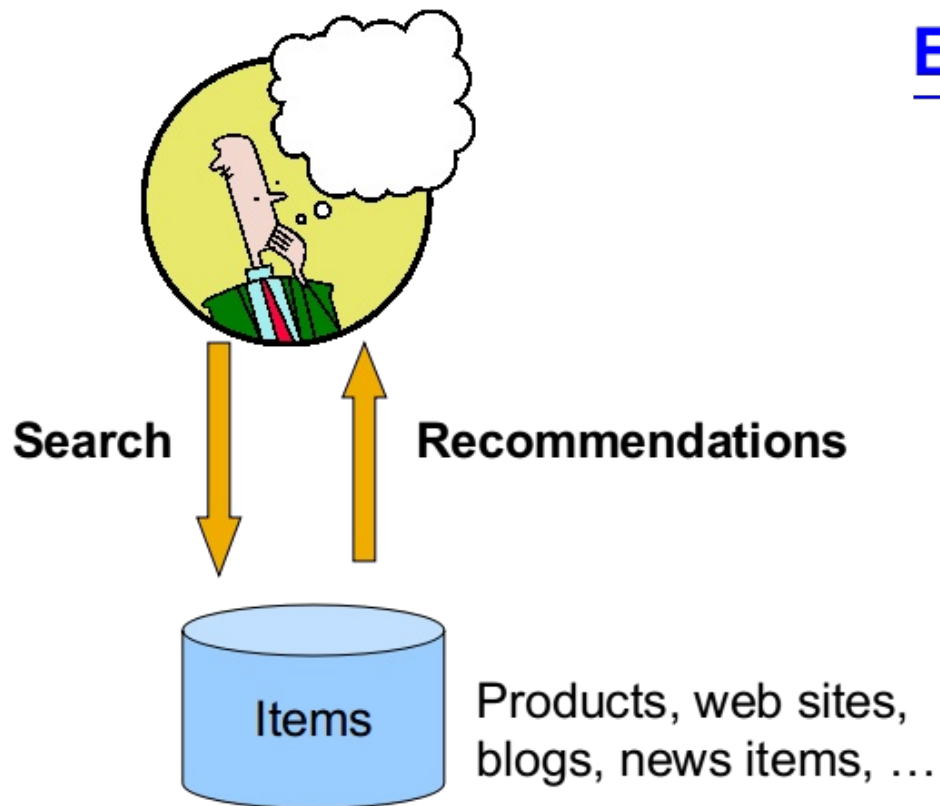
Information overload refers to the difficulty a person can have understanding an issue and making decisions that can be caused by the **presence of too much information**.

Information overload occurs when the amount of input to a system exceeds its processing capacity. Decision makers have fairly limited *cognitive processing capacity*. Consequently, when information overload occurs, it is likely that a **reduction in decision quality** will occur.

— Bertram Myron Gross (1964)

There are various solutions in computer science that can be used to mitigate information overload,

- Search
- Recommendation



Examples:

amazon.com.



StumbleUpon



del.icio.us



movie lens
helping you find the *right* movies

last.fm
the social music revolution

Google
News

You Tube


XBOX
LIVE

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

The basic idea of *Recommender Systems* is to utilize various sources of data to infer customer **interests**.

Recommendation analysis is often based on the previous **interaction** between users and items, because *past interests and proclivities are often good indicators of future choices*.

- Demographic-based
- Content-based
- **Collaborative Filtering**
 - Memory-based (kNN approaches)
 - Model-based (latent factor models)
- Hybrid



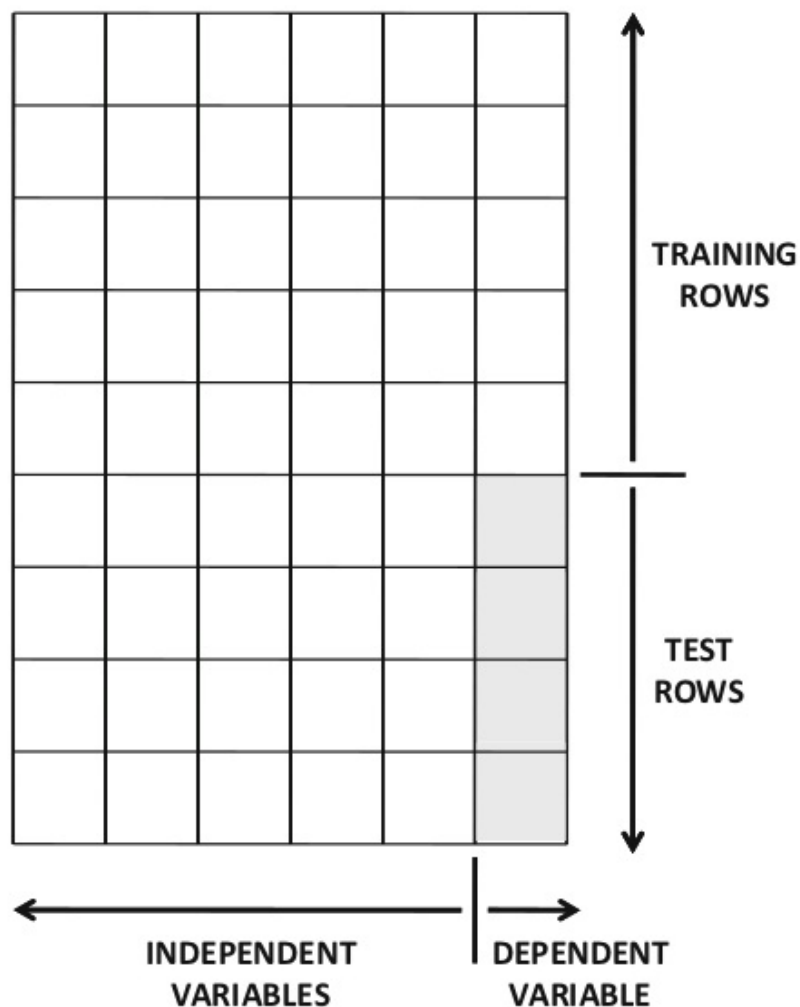
Explicit feedback:

continuous / interval-based: $\{-2, -1, 0, 1, 2\}$, forced choice rating system: $\{-2, -1, 1, 2\}$, ordinal ratings: {Strongly Disagree, Disagree, Neutral, Agree, Strongly Agree}, binary ratings: {like, dislike, unspecified}

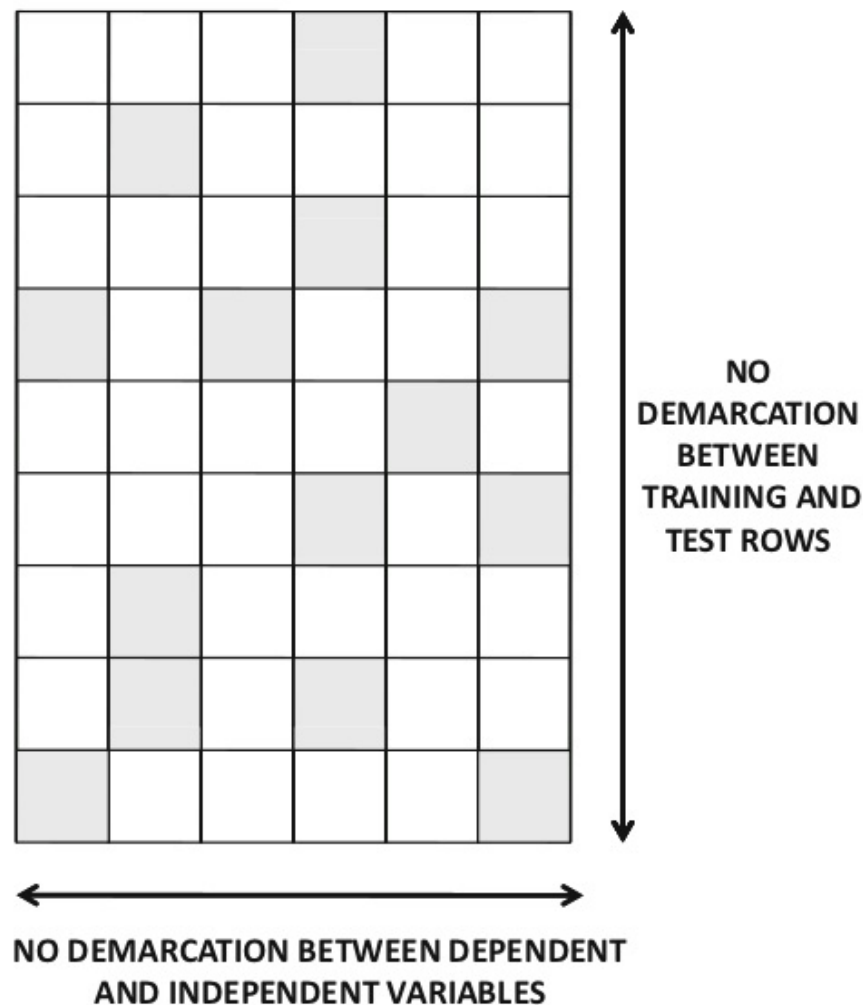
Implicit feedback:

buying, browsing, viewing, clicking behavior

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40



(a) Classification



(b) Collaborative filtering

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

$R, R_{\text{train}}, R_{\text{test}}, \hat{R}$: the set of all ratings, the training set, the test set, and the set of predicted ratings.

U : the set of all users, u and v denotes users.

I : the set of all items, i and j denotes items.

I_u : the set of all items rated by user u .

I_{uv} : the set of all items rated by both users u and v .

r_{ui} : the true rating of user u for item i .

\hat{r}_{ui} : the estimated rating of user u for item i .

μ, μ_u, μ_i : the mean of all ratings, the mean of all ratings given by user u , the mean of all ratings to item i .

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

Mean Absolute Error (MAE)

$$\text{MAE} = \frac{1}{|\hat{R}|} \sum_{\hat{r}_{ui} \in \hat{R}} |r_{ui} - \hat{r}_{ui}|$$

Root Mean Squared Error (RMSE)

$$\text{RMSE} = \sqrt{\frac{1}{|\hat{R}|} \sum_{\hat{r}_{ui} \in \hat{R}} (r_{ui} - \hat{r}_{ui})^2}$$

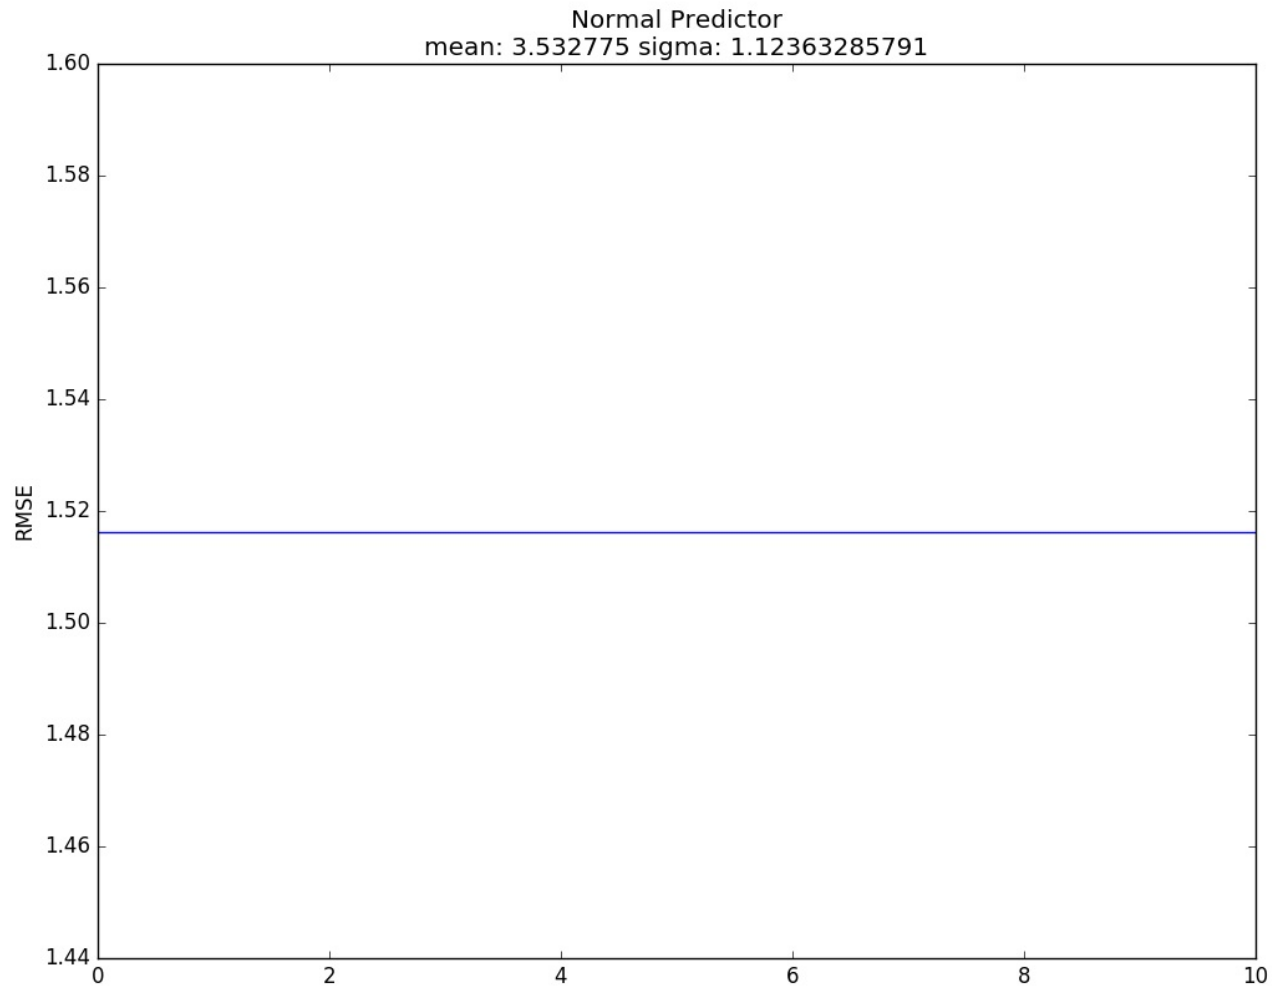
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

The prediction \hat{r}_{ui} is generated from a normal distribution $N(\hat{\mu}, \hat{\sigma}^2)$, where $\hat{\mu}$ and $\hat{\sigma}$ are estimated from the training data using Maximum Likelihood Estimation:

$$\hat{\mu} = \frac{1}{|R_{train}|} \sum_{r_{ui} \in R_{train}} r_{ui}$$

$$\hat{\sigma} = \sqrt{\sum_{r_{ui} \in R_{train}} \frac{(r_{ui} - \hat{\mu})^2}{|R_{train}|}}$$

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

Typical CF data exhibit large user and item biases—i.e., systematic tendencies for some users to give higher ratings than others, and for some items to receive higher ratings than others.

We will encapsulate those effects, which do not involve user-item interaction, within the baseline predictors (also known as biases).

$$\hat{r}_{ui} = b_{ui} = \mu + b_u + b_i$$


μ is the overall average rating. The parameters b_u and b_i indicate the observed deviations of user u and item i , respectively, from the average. b_{ui} is the baseline rating of user u for item i .

$$\min_{b_u, b_i} \sum_{u \in U, v \in V} (r_{ui} - (\mu + b_u + b_i))^2 + \lambda_u \sum_u b_u^2 + \lambda_i \sum_i b_i^2$$

gradient descent

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

ALS Baseline estimate
{'reg_u': 5, 'n_epochs': 20, 'reg_i': 5, 'method': 'als', 'rmse': 0.94067412640646286}



- Consider user x
- Find set N of other users whose ratings are “**similar**” to x ’s ratings
- Estimate x ’s ratings based on ratings of users in N



$N_i^k(u)$: the k nearest neighbors of user u that have rated item i .

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

$$\text{CosineSim}(u, v) = \frac{\sum_{i \in I_{uv}} r_{ui} \cdot r_{vi}}{\sqrt{\sum_{i \in I_{uv}} r_{ui}^2} \cdot \sqrt{\sum_{i \in I_{uv}} r_{vi}^2}}$$

$$\text{MeanSquaredDifference}(u, v) = \frac{1}{|I_{uv}|} \cdot \sum_{i \in I_{uv}} (r_{ui} - r_{vi})^2$$

$$\text{MeanSquaredDifferenceSim}(u, v) = \frac{1}{\text{msd}(u, v) + 1}$$

$$\text{PearsonCorrelation}(u, v) = \frac{\sum_{i \in I_{uv}} (r_{ui} - \mu_u) \cdot (r_{vi} - \mu_v)}{\sqrt{\sum_{i \in I_{uv}} (r_{ui} - \mu_u)^2} \cdot \sqrt{\sum_{i \in I_{uv}} (r_{vi} - \mu_v)^2}}$$

$$\text{PearsonBaselineSim}(u, v) = \hat{\rho}_{uv} = \frac{\sum_{i \in I_{uv}} (r_{ui} - b_{ui}) \cdot (r_{vi} - b_{vi})}{\sqrt{\sum_{i \in I_{uv}} (r_{ui} - b_{ui})^2} \cdot \sqrt{\sum_{i \in I_{uv}} (r_{vi} - b_{vi})^2}}$$

$$\text{PearsonBaselineShrunkSim}(u, v) = \frac{|I_{uv}| - 1}{|I_{uv}| - 1 + \text{shrinkage}} \cdot \hat{\rho}_{uv}$$

The shrinkage parameter helps to avoid overfitting when only few ratings are available.

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

$$\text{KNNBasic} \quad \hat{r}_{ui} = \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot r_{vi}}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$

$$\text{KNNWithMeans} \quad \hat{r}_{ui} = \mu_u + \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot (r_{vi} - \mu_v)}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$

$$\text{KNNWithZscore} \quad \hat{r}_{ui} = \mu_u + \sigma_u \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot z_{vi}}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$

$$z_{vi} = \frac{r_{vi} - \mu_v}{\sigma_v}$$

$$\sigma_v = \sqrt{\frac{\sum_{i \in I_v} (r_{vi} - \mu_v)^2}{|I_v| - 1}}$$

$$\text{KNNBaseline} \quad \hat{r}_{ui} = b_{ui} + \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot (r_{vi} - b_{vi})}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$

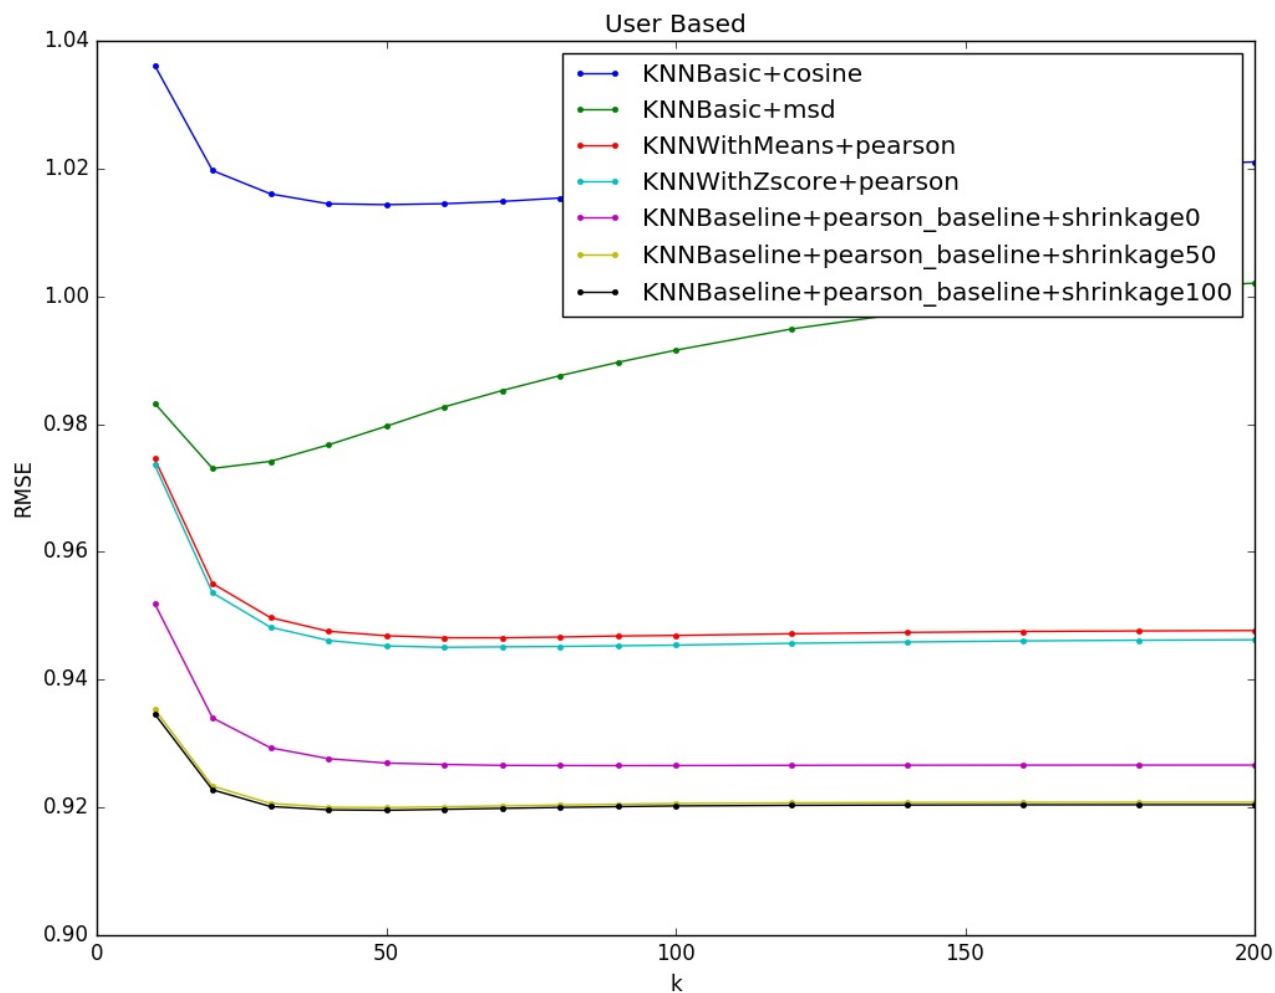
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40


Table 2.1: User-user similarity computation between user 3 and other users

Item-Id \Rightarrow	1	2	3	4	5	6	Mean Rating	Cosine($i, 3$) (user-user)	Pearson($i, 3$) (user-user)
User-Id \Downarrow									
1	7	6	7	4	5	4	5.5	0.956	0.894
2	6	7	?	4	3	4	4.8	0.981	0.939
3	?	3	3	1	1	?	2	1.0	1.0
4	1	2	2	3	3	4	2.5	0.789	-1.0
5	1	?	1	2	3	3	2	0.645	-0.817

$$\hat{r}_{31} = \frac{7 * 0.894 + 6 * 0.939}{0.894 + 0.939} \approx 6.49$$

$$\hat{r}_{36} = \frac{4 * 0.894 + 4 * 0.939}{0.894 + 0.939} = 4$$





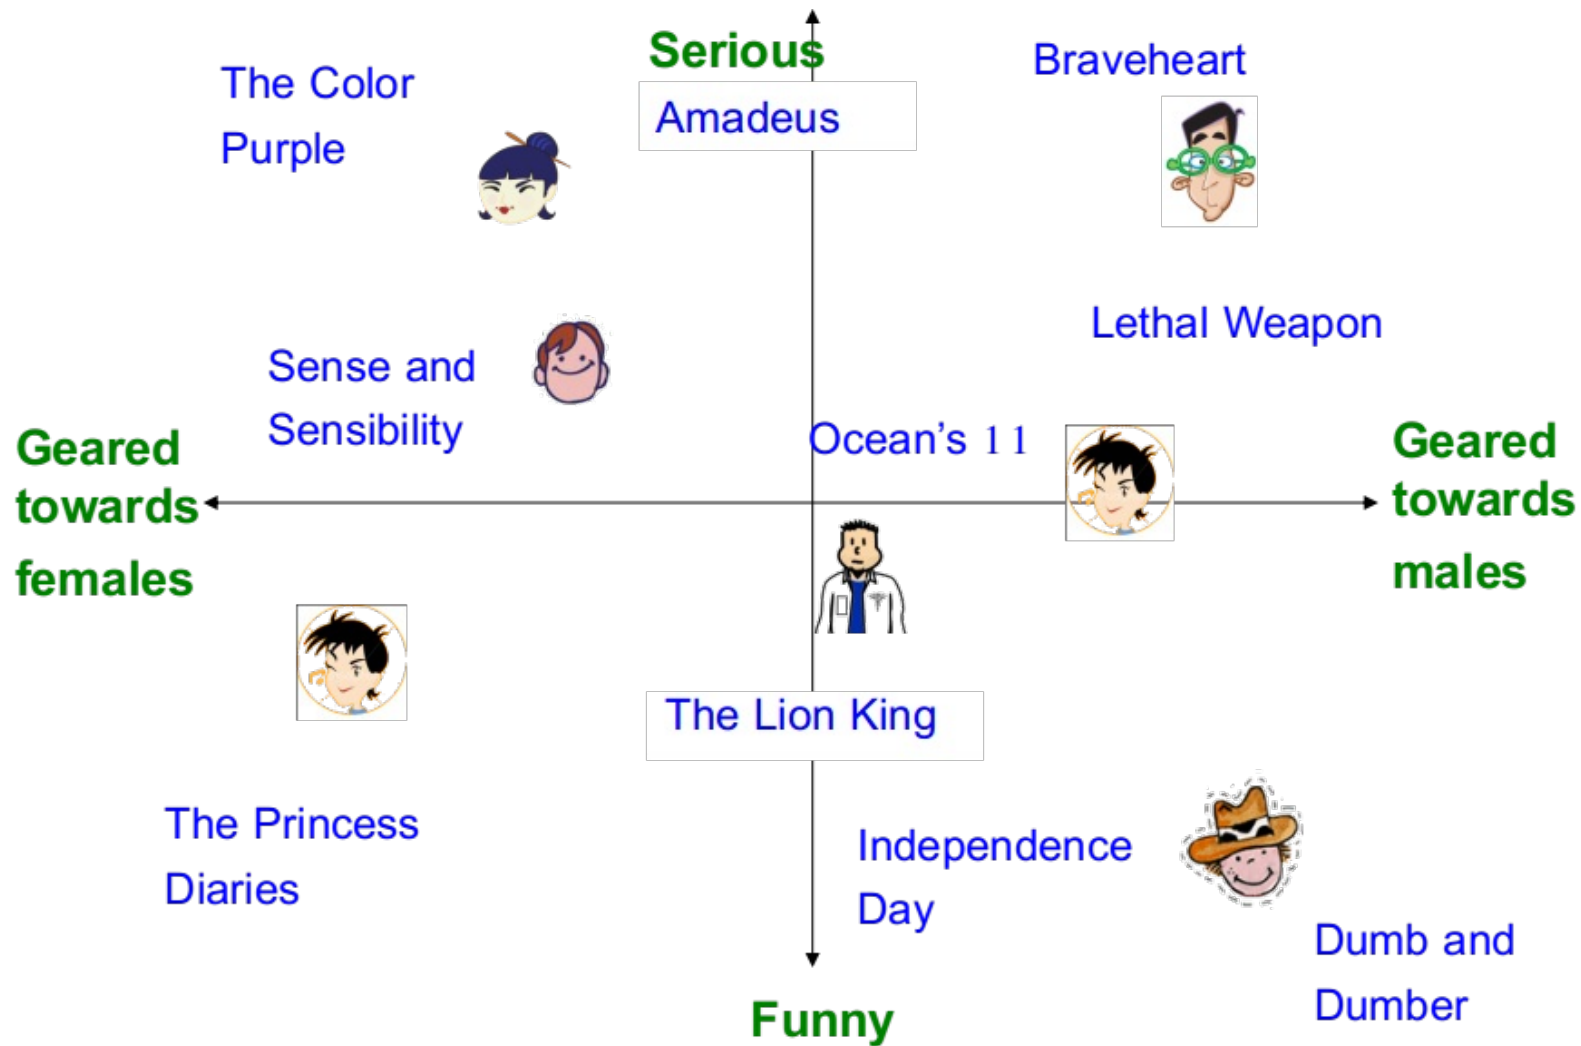
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

Pairwise similarities are hard to robustly compute in sparse rating matrices, dimensionality reduction provides a dense low-dimensional representation in terms of latent factors. Even when two users have very few items rated in common, a distance can be computed between their low-dimensional latent vectors.

The similarity computations in the reduced representation are more robust because the new low-dimensional vector is fully specified. Furthermore, the similarity computations are more efficient because of the low dimensionality of the latent representation.

The main problem of using such an approach is that the ratings matrix is **not fully specified**. As a result, this factorization is undefined.

Nevertheless, it is possible to recast the formulation as an optimization problem, in which the squared error of factorization is optimized only over the **observed** entries of the ratings matrix.



Any $m \times n$ matrix R of rank $k \ll \min \{m, n\}$ can always be expressed in the following product form of rank- k factors:

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

Let the set of all user-item pairs (i, j) , which are observed in R , be denoted by S .

$$\begin{aligned} \text{Minimize } J &= \frac{1}{2} \sum_{(i,j) \in S} e_{ij}^2 + \frac{\lambda}{2} \sum_{i=1}^m \sum_{s=1}^k u_{is}^2 + \frac{\lambda}{2} \sum_{j=1}^n \sum_{s=1}^k v_{js}^2 \\ &= \frac{1}{2} \sum_{(i,j) \in S} \left(r_{ij} - \sum_{s=1}^k u_{is} \cdot v_{js} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^m \sum_{s=1}^k u_{is}^2 + \frac{\lambda}{2} \sum_{j=1}^n \sum_{s=1}^k v_{js}^2 \\ &\quad \text{subject to: No constraints on } U \text{ and } V \end{aligned}$$

$$\begin{aligned} \frac{\partial J}{\partial u_{iq}} &= \sum_{j:(i,j) \in S} \left(r_{ij} - \sum_{s=1}^k u_{is} \cdot v_{js} \right) (-v_{jq}) + \lambda u_{iq} = \sum_{j:(i,j) \in S} (e_{ij})(-v_{jq}) + \lambda u_{iq} \\ &\quad \forall j \in \{1 \dots m\}, q \in \{1 \dots k\} \end{aligned}$$

$$\begin{aligned} \frac{\partial J}{\partial v_{jq}} &= \sum_{i:(i,j) \in S} \left(r_{ij} - \sum_{s=1}^k u_{is} \cdot v_{js} \right) (-u_{iq}) + \lambda v_{jq} = \sum_{i:(i,j) \in S} (e_{ij})(-u_{iq}) + \lambda v_{jq} \\ &\quad j \in \{1 \dots n\}, q \in \{1 \dots k\} \\ u_{iq} &\Leftarrow u_{iq} + \alpha \frac{\partial J}{\partial u_{iq}}, v_{jq} \Leftarrow v_{jq} + \alpha \frac{\partial J}{\partial v_{jq}} \end{aligned}$$

$$\begin{aligned} \text{Minimize } J &= \frac{1}{2} \sum_{(i,j) \in S} e_{ij}^2 + \frac{\lambda}{2} \sum_{i=1}^m \sum_{s=1}^k u_{is}^2 + \frac{\lambda}{2} \sum_{j=1}^n \sum_{s=1}^k v_{js}^2 \\ &= \frac{1}{2} \sum_{(i,j) \in S} \left(r_{ij} - \sum_{s=1}^k u_{is} \cdot v_{js} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^m \sum_{s=1}^k u_{is}^2 + \frac{\lambda}{2} \sum_{j=1}^n \sum_{s=1}^k v_{js}^2 \end{aligned}$$

subject to:

Columns of U are mutually orthogonal

Columns of V are mutually orthogonal

Non negative Matrix Factorization

subject to:

$$U \geq 0$$

$$V \geq 0$$

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

■ Training data

- 100 million ratings, 480,000 users, 17,770 movies
- 6 years of data: 2000-2005

■ Test data

- Last few ratings of each user (2.8 million)
- **Evaluation criterion:** Root Mean Square Error (RMSE)


$$= \frac{1}{|R|} \sqrt{\sum_{(i,x) \in R} (\hat{r}_{xi} - r_{xi})^2}$$

- **Netflix's system RMSE: 0.9514**


■ Competition

- 2,700+ teams
- **\$1 million** prize for 10% improvement on Netflix

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40


user bias



movie bias



user-movie interaction



Baseline predictor

- Separates users and movies
- Benefits from insights into user's behavior
- Among the main practical contributions of the competition

User-Movie interaction

- Characterizes the matching between users and movies
- Attracts most research in the field
- Benefits from algorithmic and mathematical innovations

- μ = overall mean rating
- b_x = bias of user x
- b_i = bias of movie i

$$r_{xi} = \underbrace{\mu}_{\substack{\text{Overall} \\ \text{mean} \\ \text{rating}}} + \underbrace{b_x}_{\substack{\text{Bias for} \\ \text{user } x}} + \underbrace{b_i}_{\substack{\text{Bias for} \\ \text{movie } i}} + \underbrace{q_i \cdot p_x}_{\substack{\text{User-Movie} \\ \text{interaction}}}$$

■ Example:

- Mean rating: $\mu = 3.7$
- You are a critical reviewer: your ratings are 1 star lower than the mean: $b_x = -1$
- Star Wars gets a mean rating of 0.5 higher than average movie: $b_i = +0.5$
- Predicted rating for you on Star Wars:
 $= 3.7 - 1 + 0.5 = 3.2$

■ Solve:

$$\min_{Q,P} \sum_{(x,i) \in R} \left(r_{xi} - (\mu + b_x + b_i + q_i p_x) \right)^2$$

goodness of fit

$$+ \left(\lambda_1 \sum_i \|q_i\|^2 + \lambda_2 \sum_x \|p_x\|^2 + \lambda_3 \sum_x \|b_x\|^2 + \lambda_4 \sum_i \|b_i\|^2 \right)$$


regularization

λ is selected via grid-search on a validation set

■ Stochastic gradient decent to find parameters



- **Note:** Both biases \mathbf{b}_x , \mathbf{b}_i as well as interactions \mathbf{q}_i , \mathbf{p}_x are treated as parameters (we estimate them)

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40



- **Sudden rise in the average movie rating (early 2004)**
 - Improvements in Netflix
 - GUI improvements
 - Meaning of rating changed
- **Movie age**
 - Users prefer new movies without any reasons
 - Older movies are just inherently better than newer ones

Y. Koren, Collaborative filtering with temporal dynamics, KDD '09



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

- **Original model:**

$$r_{xi} = \mu + b_x + b_i + q_i \cdot p_x$$


- **Add time dependence to biases:**

$$r_{xi} = \mu + b_x(t) + b_i(t) + q_i \cdot p_x$$

- Make parameters b_x and b_i to depend on time
- (1) Parameterize time-dependence by linear trends
- (2) Each bin corresponds to 10 consecutive weeks

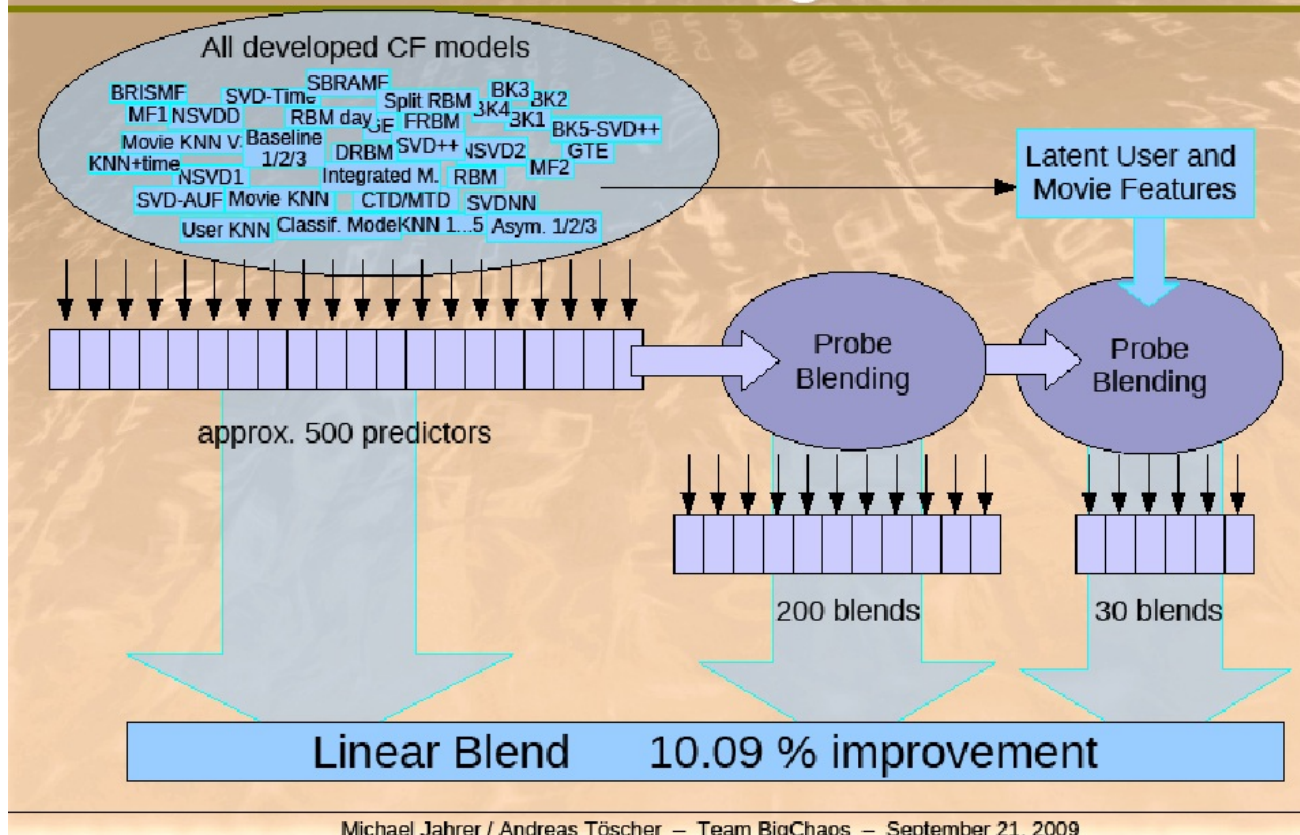
- **Add temporal dependence to factors**

- $p_x(t)$... user preference vector on day t



The big picture

Solution of BellKor's Pragmatic Chaos



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

- **The winner of the Netflix Challenge!**

- **Multi-scale modeling of the data:**

Combine top level, “regional” modeling of the data, with a refined, local view:

- **Global:**

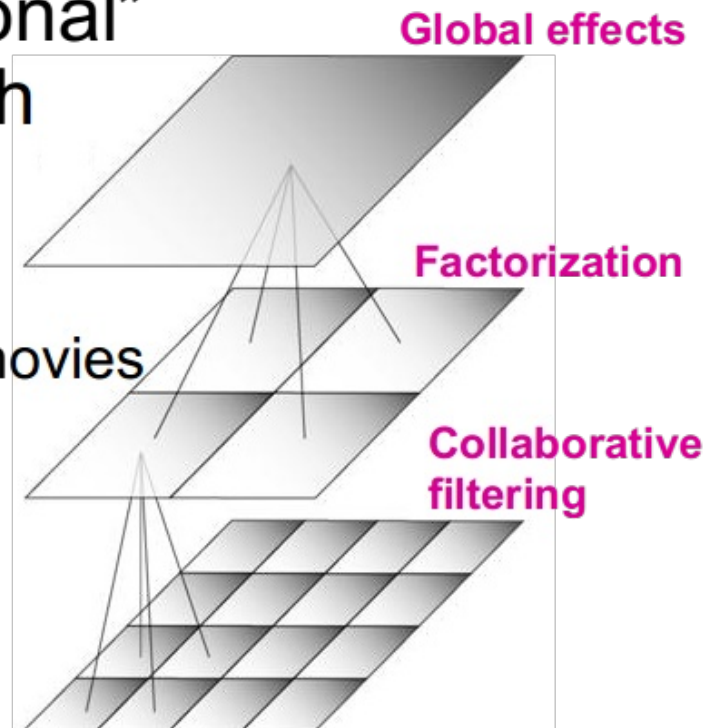
- Overall deviations of users/movies

- **Factorization:**

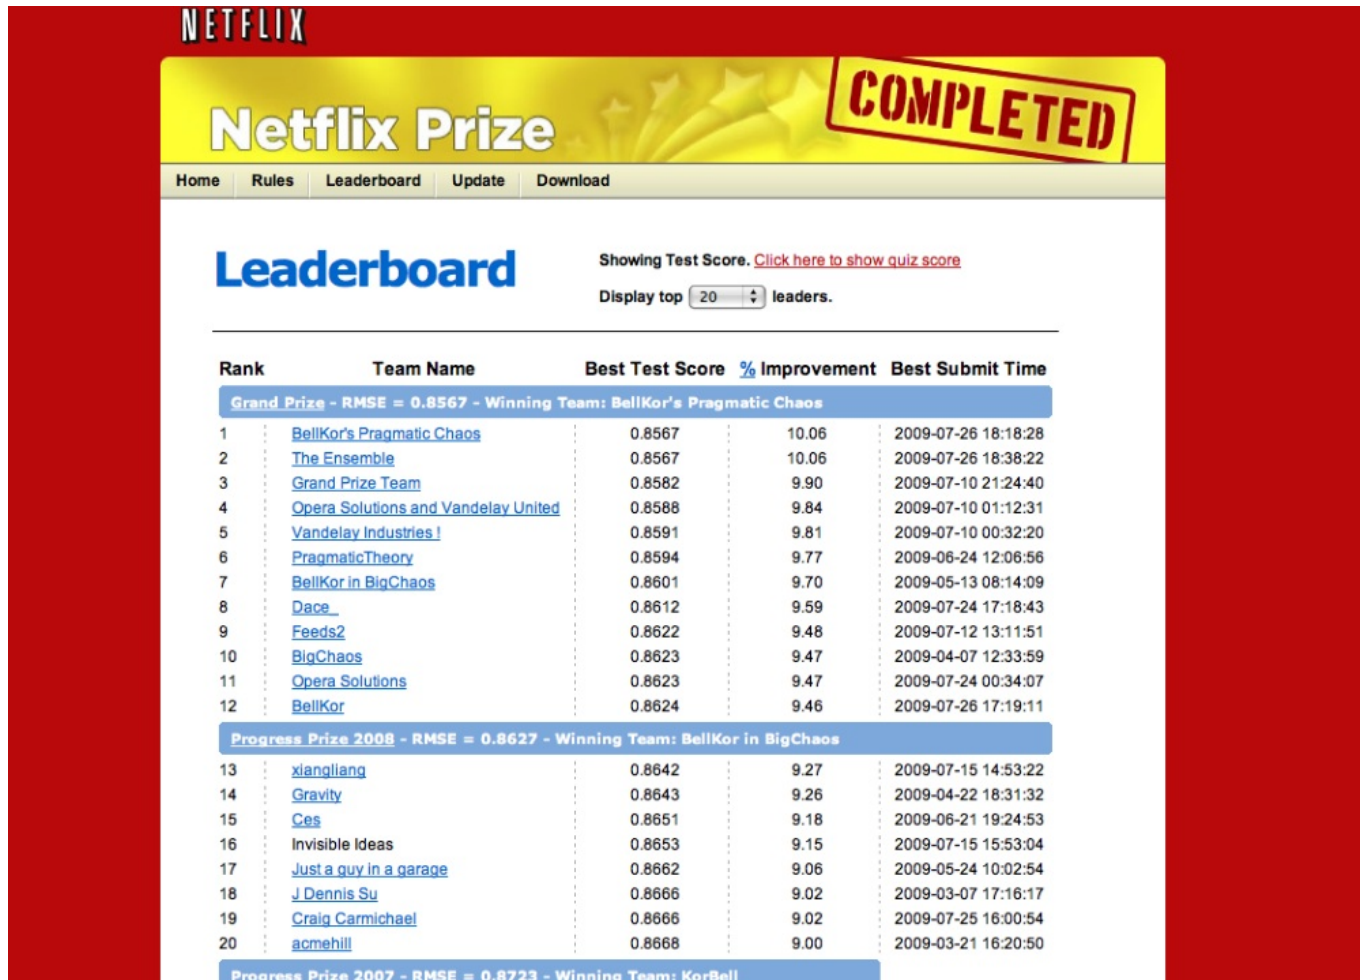
- Addressing “regional” effects

- **Collaborative filtering:**

- Extract local patterns



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40



The screenshot shows the Netflix Prize Leaderboard interface. At the top, the Netflix logo is on the left, and a large yellow banner with the text "Netflix Prize" and a red "COMPLETED" stamp is on the right. Below the banner is a navigation bar with links: Home, Rules, Leaderboard, Update, and Download. The main heading "Leaderboard" is in large blue text. To its right, it says "Showing Test Score. [Click here to show quiz score](#)". Below this, a "Display top" dropdown menu is set to "20" and followed by the text "leaders.". The leaderboard table has five columns: Rank, Team Name, Best Test Score, % Improvement, and Best Submit Time. It is divided into three sections: Grand Prize (Ranks 1-12), Progress Prize 2008 (Ranks 13-19), and Progress Prize 2007 (Rank 20). Each section header includes the RMSE value and the winning team name.

Rank	Team Name	Best Test Score	% Improvement	Best Submit Time
Grand Prize - RMSE = 0.8567 - Winning Team: BellKor's Pragmatic Chaos				
1	BellKor's Pragmatic Chaos	0.8567	10.06	2009-07-26 18:18:28
2	The Ensemble	0.8567	10.06	2009-07-26 18:38:22
3	Grand Prize Team	0.8582	9.90	2009-07-10 21:24:40
4	Opera Solutions and Vandelay United	0.8588	9.84	2009-07-10 01:12:31
5	Vandelay Industries !	0.8591	9.81	2009-07-10 00:32:20
6	PragmaticTheory	0.8594	9.77	2009-06-24 12:06:56
7	BellKor in BigChaos	0.8601	9.70	2009-05-13 08:14:09
8	Dace	0.8612	9.59	2009-07-24 17:18:43
9	Feeds2	0.8622	9.48	2009-07-12 13:11:51
10	BigChaos	0.8623	9.47	2009-04-07 12:33:59
11	Opera Solutions	0.8623	9.47	2009-07-24 00:34:07
12	BellKor	0.8624	9.46	2009-07-26 17:19:11
Progress Prize 2008 - RMSE = 0.8627 - Winning Team: BellKor in BigChaos				
13	xiangliang	0.8642	9.27	2009-07-15 14:53:22
14	Gravity	0.8643	9.26	2009-04-22 18:31:32
15	Ces	0.8651	9.18	2009-06-21 19:24:53
16	Invisible Ideas	0.8653	9.15	2009-07-15 15:53:04
17	Just a guy in a garage	0.8662	9.06	2009-05-24 10:02:54
18	J Dennis Su	0.8666	9.02	2009-03-07 17:16:17
19	Craig Carmichael	0.8666	9.02	2009-07-25 16:00:54
20	acmehill	0.8668	9.00	2009-03-21 16:20:50
Progress Prize 2007 - RMSE = 0.8723 - Winning Team: KorBell				



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40

Leskovec, J., Rajaraman, A., & Ullman, J. D. (2014). *Mining of massive datasets*. Cambridge university press. <http://www.mmids.org>

Abu-Mostafa, Y. S., Magdon-Ismail, M., & Lin, H. T. (2012). *Learning from data* (Vol. 4). New York, NY, USA:: AMLBook.

Kantor, P. B. (2015). *Recommender systems handbook*. F. Ricci, L. Rokach, & B. Shapira (Eds.). Berlin, Germany:: Springer.

Aggarwal, C. C. (2016). *Recommender systems*. Springer International Publishing.

<http://surprise.readthedocs.io/en/stable/index.html>